




$$u_t = \text{normalize}(u_L + \tan(\theta_{max}) \Delta u_\perp), \quad \Delta u_\perp = \Delta u_t - (\Delta u_t \cdot u_L) u_L$$

Lead uses the same sanitized target-state observation as the residual policy.

Output is a desired velocity DIRECTION only; the environment enforces all physical limits.